

BPF Tales




**Or: Why Did I Recompile the
Kernel Just to Average Numbers?**


We need some context

The data deluge


- 44ZB of stored data 2020
- 53% of companies use Big Data analytics
- \$77 billion worth by 2023




So, this is a datacenter



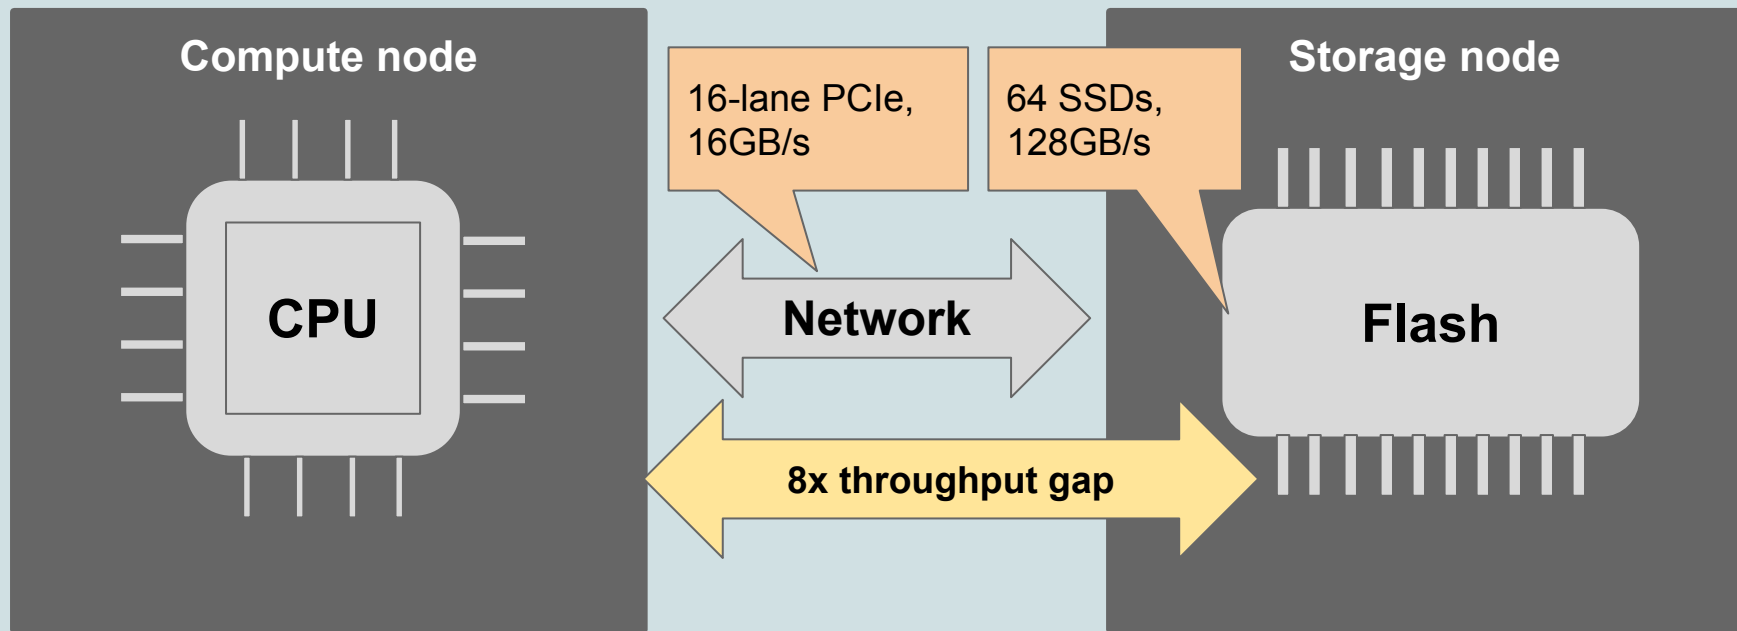
So, this is a datacenter



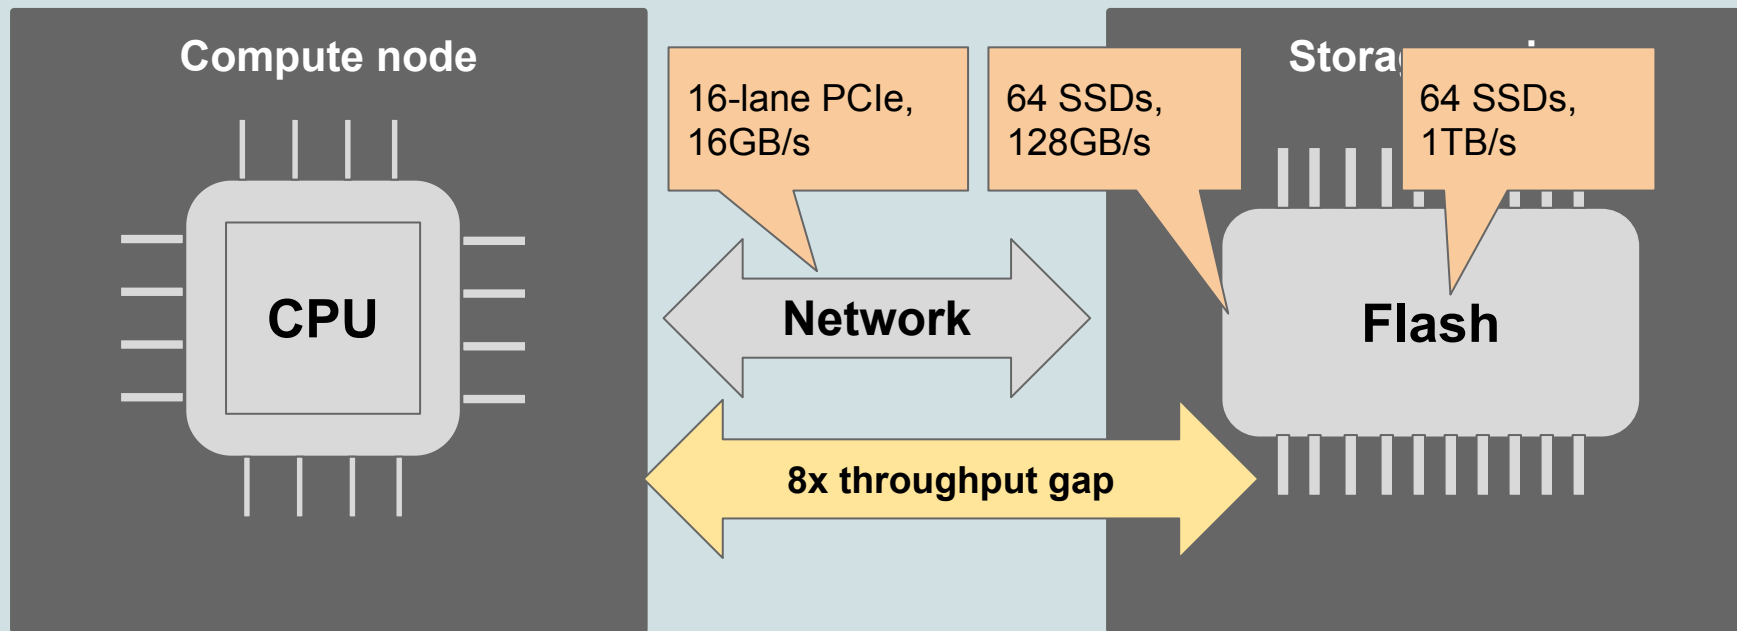
So, this is a datacenter



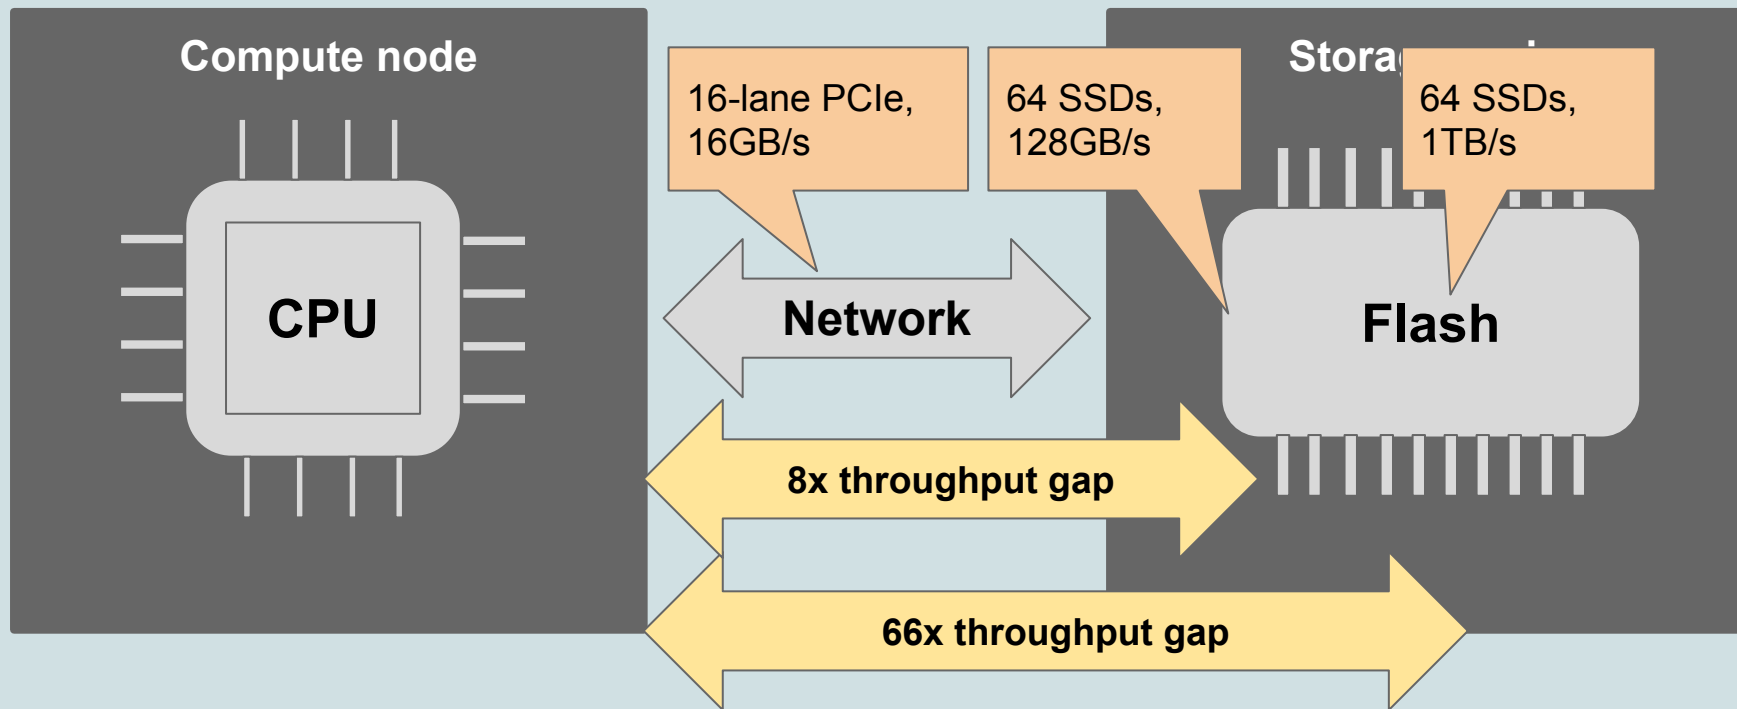
So, this is a datacenter



So, this is a datacenter



So, this is a datacenter



What is happening?

- The data deluge is only getting worse
- Storage is becoming TOO fast
- Data movement is a bottleneck

**So this is what
I worked on**

**For
almost
a year**



How to fix this?


- Reduce data movement
- Compute on storage
- Optimize data transfers



Programmability!

How to fix this?

- Reduce data movement
- Compute on storage
- Optimize data transfers



Programmability!




Who am I


- Giulia Frascaria
- Vrije Universiteit, Amsterdam
- PhD candidate
- atLarge research group

Ok, let's start...


Ok, let's average numbers




Ok, let's average numbers




Ok, let's compute the average




Ok, let's compute the average




Ok, let's compute the average




Ok, let's compute the average




Ok, let's compute the average



Ok, let's compute the average




Ok, let's compute the average




This doesn't make sense

- Huge data transfer
- High latency
- Unnecessary network congestion
- Bottleneck on storage


Let's move it to the storage!




Let's move it to the storage!




Let's move it to the storage!



Let's move it to the storage!



Let's move it to the storage!



This makes sense!

- Only transfer result
- Reduced network congestion
- Use storage throughput

How to fix this?

Datacenter requirements

- **Multitenancy**
- **Isolation and security**
- **High performance**
- **Low deployment cost**

And my requirement

- Multitenancy
- Isolation and security
- High performance
- Low deployment cost
- Reduce data movement


What are the options?

- Python, Rust, Go...
- ASIC, FPGAs
- What about kernel space?



What are the options?

- Python, Rust, Go...
- ASIC, FPGAs
- What about kernel space?
- eBPF is included



eBPF 101

- Javascript for the kernel
- Programmable
- In-kernel VM
- Efficient
- Safe and formally verified

eBPF for networking

- Packet inspection
- Modify packets
- Drop packets

eBPF for networking

- ~~Packet inspection~~ READ
- ~~Modify packets~~ WRITE
- ~~Drop packets~~ DROP

How can I do it
on storage?



Warning:
kernel code ahead

The read() path

+ 94.378 us

0.057 us

0.538 us

0.065 us

0.494 us

0.902 us

0.058 us

0.068 us

2.216 us

2.655 us

0.064 us

0.503 us


0.365 us

```
}
Sys_read() {
    __fdget_pos() {
        __fget_light();
    }
    vfs_read() {
        rw_verify_area() {
            security_file_permission() {
                apparmor_file_permission() {
                    common_file_perm() {
                        aa_file_perm();
                    }
                }
            }
            __fsnotify_parent();
            fsnotify();
        }
    }
    __vfs_read() {
        ext4_file_read_iter() {
            generic_file_read_iter() {
                _cond_resched() {
                    rcu_all_qs();
                }
            }
            pagecache_get_page() {
                find_get_entry();
            }
        }
    }
}
```

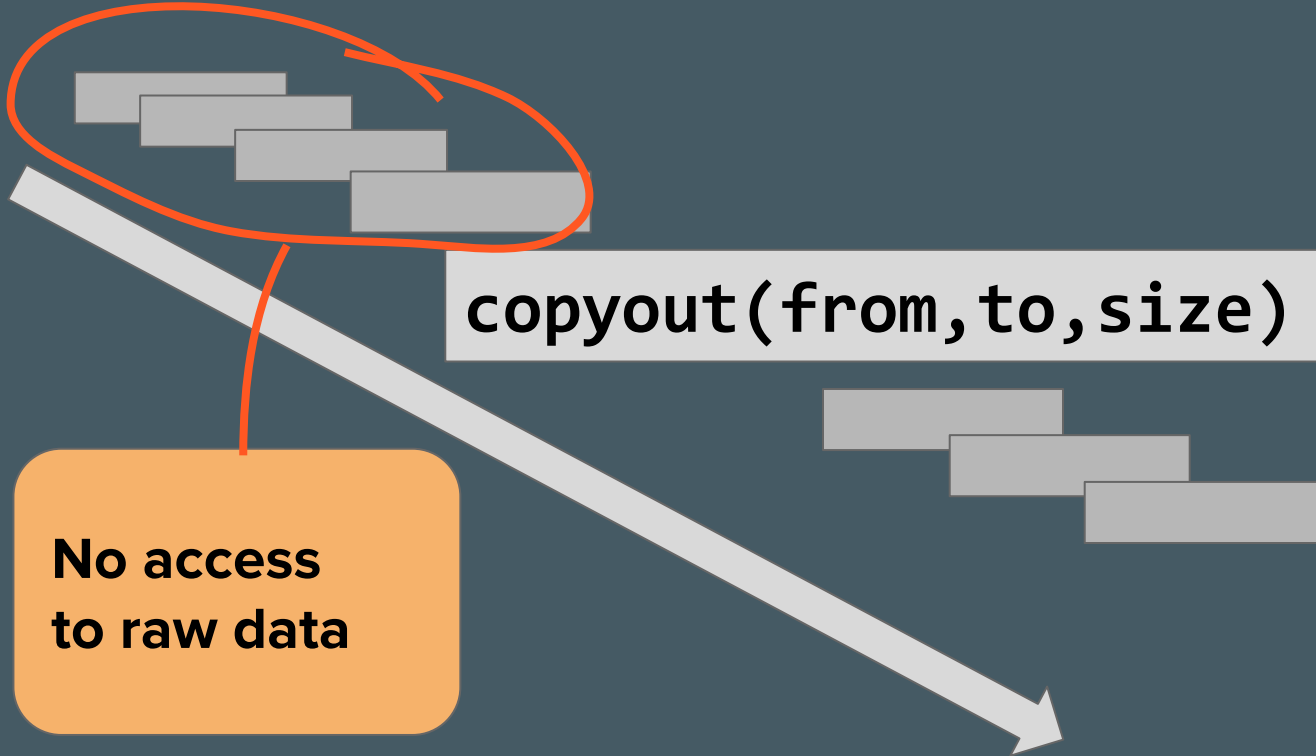
What am I looking for?

- Identify user and file
- Access to the data
- Execute custom code

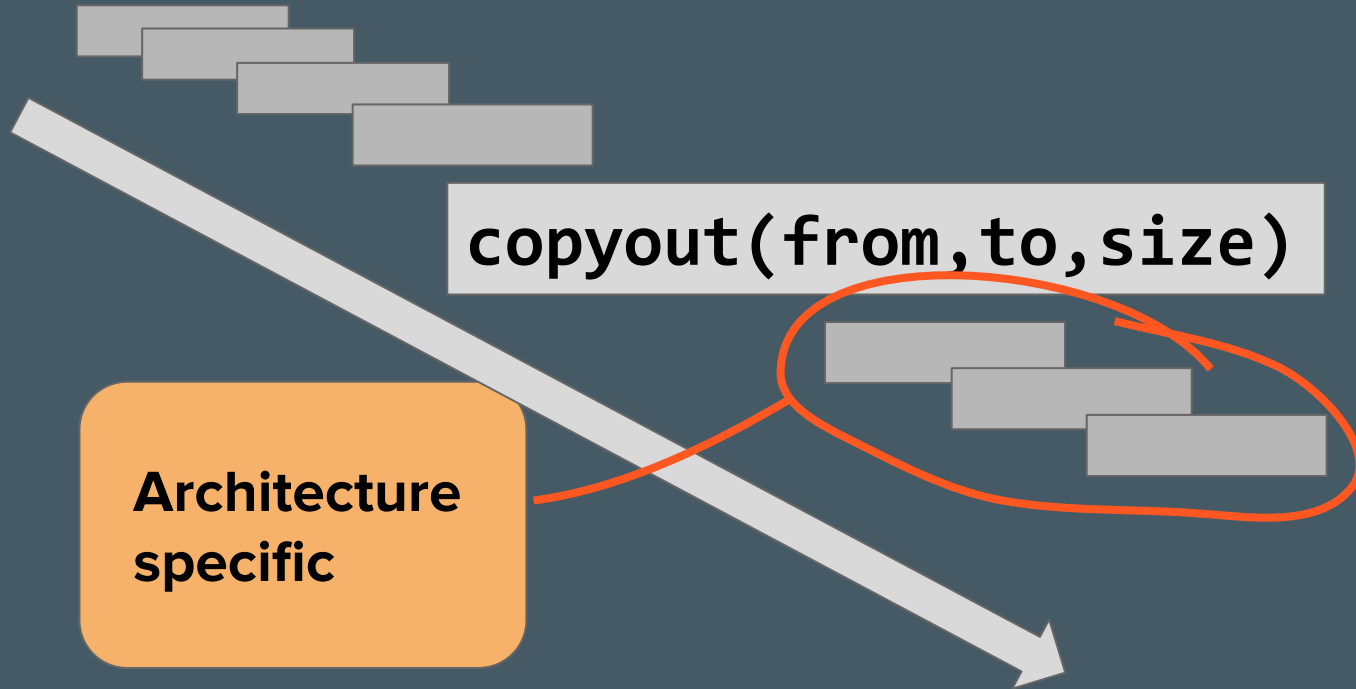
Somewhere in the read path...




Somewhere in the read path...



Somewhere in the read path...




Somewhere in the read path...




So, all according to plan?


This is the prototype



This is the prototype




This is the prototype




Why Filter-Reduce?

- Keep a standard interface
- Offer users some flexibility
- MapReduce
- Relational data processing
- Maximum data reduction


So actually this happens




Let's move it to the storage!



Let's move it to the storage!



Let's move it to the storage!



Thanks to eBPF:

- Reduce network congestion
- Avoid kernel to user copy
- Users do not read data
- Only result is shared

eBPF runs in a kernel sandbox

So, all according to plan?

Well you read the title

First of all, where to hook the bpf extensions

- Copyout is not an exported symbol
- Necessary to use modified Linux kernel

Second, what about data access?

- Even read-only, you need a helper function
- Adds overhead
- Direct pointer access would be more efficient

Third, how much can I process?

- No dynamic allocation
- Stack limited to 512 bytes
- Need to iterate in batches
- eBPF instruction number limit

Well at least it computes the average, right?

- Allow to use helpers to convert char to int
- Floating point division not supported in kernel

So, to summarize

- Conceptually same as networking
- Achieves data movement reduction
- Safe, isolated execution

But we need more support for I/O!

Where are we headed

eBPF is not slowing down

- Tracing: BCC, bpftrace...
- Networking: Cilium...
- Security: KRSI, Falco...

Is programmability the next step?

We're not alone, will you join?

DEVELOPMENT / CONTRIBUTED

How io_uring and eBPF Will Revolutionize Programming in Linux

21 Apr 2020 8:49am, by [Glauber Costa](#)

DEVOPS

eBPF - Rethinking the Linux Kernel

LIKE

DISCUSS



DEVELOPMENT / LINUX

How eBPF Turns Linux into a Programmable Kernel

8 Oct 2020 6:00am, by [Joab Jackson](#)

https://thenewstack.io/how-io_uring-and-ebpf-will-revolutionize-programming-in-linux/

<https://thenewstack.io/how-ebpf-turns-linux-into-a-programmable-kernel/>

<https://www.infoq.com/presentations/facebook-google-bpf-linux-kernel/>